

# Object-oriented wrappers for the Linux kernel



D. Janakiram, Ashok Gunnam<sup>\*,†</sup>, N. Suneetha, Vineet Rajani  
and K. Vinay Kumar Reddy

*Distributed & Object Systems Lab, IIT Madras, Chennai, India*

## SUMMARY

Linux is an open-source operating system, which has increased in its popularity and size since its birth. Various studies have been conducted in literature on the evolution of the Linux kernel, which have shown that there are considerable maintenance problems arising out of the coupling issues in the Linux kernel and this may hamper the evolution of the kernel in future. We propose an object-oriented (OO) wrapper-based approach to Linux kernel to provide OO abstractions to external modules. As the major growth of the size of the Linux kernel is in device drivers, our approach provides substantial benefits in terms of developing the device drivers in C++, although the kernel is in C. Providing reusability and extensibility features to device drivers improves the maintainability of the kernel. The OO wrappers provide several benefits to module developers in terms of understandability, development ease, support for OO modules, etc. The design and implementation of C++ wrappers for Linux kernel and the performance of a device driver re-engineered in C++ are presented in this paper. Copyright © 2008 John Wiley & Sons, Ltd.

*Received 13 May 2007; Revised 3 January 2008; Accepted 13 January 2008*

KEY WORDS: Linux kernel; API; object orientation; wrappers; device drivers

## 1. INTRODUCTION

Linux is an open-source operating system (OS) developed using the C language. It is generally argued that performance is the most important quality factor for an OS. Consequently, C and assembly language are used for their development and the programs are carefully optimized for

<sup>\*</sup>Correspondence to: Ashok Gunnam, Distributed & Object Systems Lab, IIT Madras, Chennai, India.

<sup>†</sup>E-mail: ashok.gunnam@gmail.com

Contract/grant sponsor: Department of Information Technology, Government of India

performance. Often, the techniques used, although they improve the performance of the OS, they result in the degradation of maintenance.

The increasing attention on performance has resulted in increasing the complexity of the Linux kernel. Various studies [1–3] have been conducted in literature on the evolution of Linux kernel, which show evidence of considerable maintenance problems arising out of coupling in the kernel and this may hamper the evolution of the kernel in future. It has also affected the understandability of the kernel due to lack of high-level abstractions.

One possible solution to the maintenance problem is to restructure the kernel using object-oriented (OO) principles and re-implement the kernel in an OO language such as C++. However, the Linux community is reluctant to adopt such an approach mainly citing performance problems with mainstream OO languages such as C++ [4]. There are attempts to use the OO principles such as virtual tables for polymorphic function dispatch rather than use C++ language for implementing the kernel. Some of these arguments are well founded and we argue in this paper that an acceptable approach would be to provide OO wrappers for the core kernel so that outside modules such as device drivers could be developed in C++.

Device drivers constitute a major part of the Linux kernel code. The device driver code lies outside the core kernel. In addition, the major evolution of the kernel is in device drivers to support the evolving hardware trends. Fifty-seven per cent of the kernel code is observed to be of device drivers [5]. These drivers are developed as kernel modules that can be linked either statically or at runtime. Hence, the kernel can be viewed as shown in Figure 1, consisting of external kernel modules (such as device drivers) and the core kernel.

One approach towards providing better maintainability for the OS would be to develop the entire kernel and hence the OS in an OO language. An OO OS has many advantages such as extensibility, reconfigurability and reusability [6,7]. Development of programs, such as device drivers, on top of the OO OS brings out the advantages of OO programming in that each hardware device can be treated as an object. C++ as a suitable OO language for the development of the OS kernel is well argued [6,8]. However, as most OSs are developed in C, it may not be a feasible solution to re-engineer the whole kernel into C++, due to factors such as portability of existing applications, human expertise, cost involved for this and so on.

An alternate and light weight solution is to provide OO wrappers for the core kernel. The underlying kernel remains the same but appears as an OO core kernel to external modules, such as device drivers. We have applied this design to Linux kernel and provided C++ wrappers for the core kernel of Linux (2.6.9).

In this paper, we present the design, implementation and the performance evaluation of the C++ wrapper-based Linux kernel. To evaluate the usefulness of C++ wrappers, some of the existing device drivers have been re-engineered into C++. The performance evaluation of a set of C++ drivers has shown that there is a performance degradation of less than 2%, which we believe is an acceptable overhead, given the advantages that are derived from the approach.




Figure 1. Linux kernel.

## 2. TECHNICAL CHALLENGES

1. Identifying abstractions: An important feature of object orientation is *abstraction*. According to OO philosophy, abstraction should be based on data, i.e. the program should be structured around data rather than functionality. However, in the procedural paradigm, programs are usually structured based on functionality (for example, using flow charts during their design). As Linux is developed using a procedural language, it is structured based on functionality [3]. For this reason, it is necessary to identify the candidate cohesive objects from the Linux kernel, which are abstracted based on data.
2. C++ support: The original Linux kernel does not provide a C++ execution environment. Some of the C++ features such as templates, exceptions, global and static objects are not supported in kernel mode programming. Moreover, some of the keywords specific to C++ such as *new* and *delete* are used as identifiers in the Linux kernel. To handle these issues, an additional support is required to compile the C++ modules that use the original kernel header files. Also, the C++ libraries required for runtime support of the C++ modules need to be provided.

## 3. IDENTIFICATION OF OBJECTS

Many approaches for identification of abstractions from procedural code have been discussed in literature [9–11]. The common principle is to identify the global data and data structures from procedural code and group them with the functions that use them. The rationale behind this kind of grouping is to structure the program based on data rather than functionality. According to this philosophy, all the procedures that use or modify the same data belong to the same module, which in this case is a *class*.

Identification of the objects and the relationships between them, with necessary callback objects are the main requirements for identifying OO abstractions for the kernel.

The following are the set of heuristics that are used to identify the objects:

1. Identify the types (data structures) that are present in the code.
2. Identify the functions that operate on these types, i.e. either formal parameters or return types. The data structures that are used inside the body of the function are used to identify the relationships among objects.
3. Group the identified type and its appropriate functions into an abstraction.
4. If a function operates on more than one type and any of them is modified inside the function, the function is grouped with the type that gets modified.
5. If a function uses more than one data structure, and if no type is modified, then assign the function to the data structure, which is the supertype of other types, i.e. composite of other data structures.

For example, *mm\_struct* is a data type that stores information related to process' memory. Some of the procedures, which use this data structure as either formal parameters or return types, are *mm\_alloc*, *exit\_mm*, *handle\_mm\_fault* and *mmaput*. Hence, these functions are grouped together to form a class. Among these functions, *handle\_mm\_fault* accesses other data structures such as

*vm\_area\_struct* in addition to *mm\_struct*. The function assignment collision is resolved as specified in step 5. As *vm\_area\_struct* is a subtype of *mm\_struct*, the function is assigned to *mm\_struct*.

### 3.1. Callback objects

The Linux kernel uses function pointers to implement callbacks, i.e. the kernel calls the function pointed to by the function pointers when certain events such as timer or new module insertion occurs. Callbacks are important functions required by most of the external modules, such as device drivers. Hence, there should be a way to provide these callbacks in the wrappers in C++.

An easy way to identify callbacks is to identify the function pointers that are declared in the kernel. To be able to provide an OO way of callbacks and still be able to use the underlying kernel callback infrastructure, we used a technique, which is a combination of template and strategy design patterns.

The kernel implements the callbacks by declaring function pointers in kernel in one of the following two ways:

1. Inside a structure (for example, *file\_operations*).
2. As parameters for a function.

In case 1, a design that is based on template method pattern [12] is used. We declare a new class that contains virtual functions, each of which represents a corresponding function pointer declared in the structure. A C interface is defined and is statically registered with the function pointers of kernel. This C interface is a kind of template method. It calls the hooks that are virtual functions of the classes declared. These virtual functions are called in a polymorphic way. A user can define a subclass of the corresponding superclass, define the virtual functions in it, and the callbacks are registered automatically.

For example, Figure 2 shows a callback object around the *file\_operations* structure.

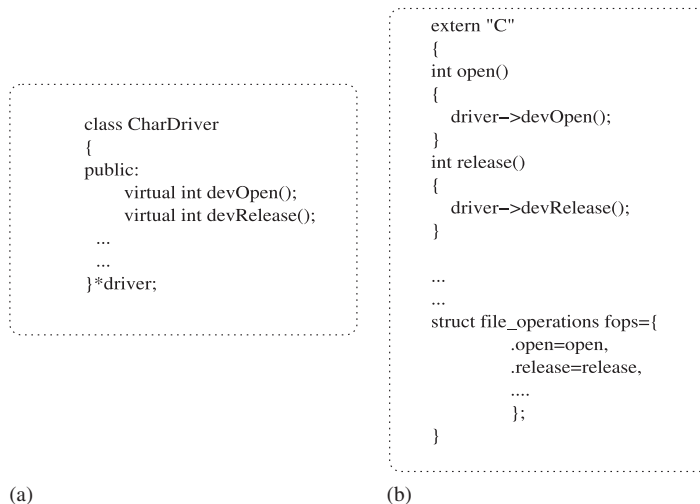


Figure 2. Abstractions for function pointers inside a structure.

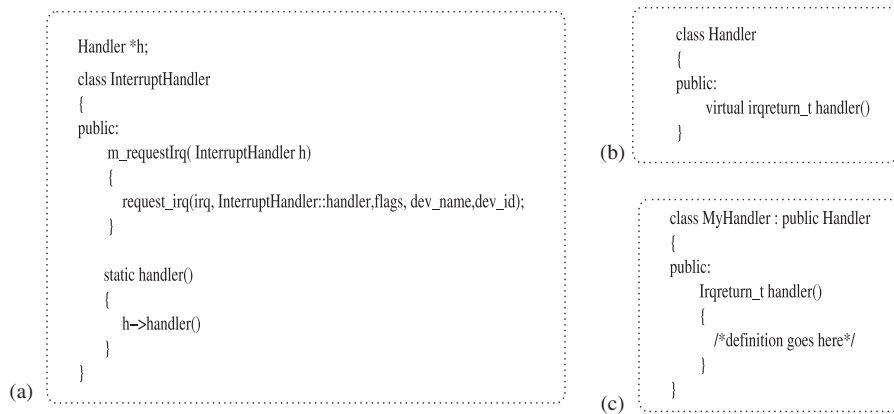


Figure 3. Abstractions for function pointers used as parameters to a function.

For case 2, a strategy pattern is used. The strategy pattern states that the algorithm be encapsulated so that the actual algorithm needed can be assigned at runtime. In case 2, where function pointers are passed as parameters to another function, it is similar to an OO strategy pattern. Here there is no need to register a C interface to the kernel statically.

For example, Figure 3 shows how interrupt handlers of the Linux kernel are provided in wrappers. The `request_irq` function of Linux kernel has a pointer to handler function, which the user defines at a later point in time. Here, the algorithm that is defined at runtime is the actual handler function, which in OO terms is the strategy. The context class, which in this case is `InterruptHandler` class, has a reference to a strategy and delegates the request to the appropriate algorithm at runtime based on the concrete strategy assigned to the reference.

## 4. DESIGN

This section explains the design of C++ wrapper-based core kernel and the design goals in accomplishing the task.

### 4.1. Design goals

1. *Structure should be identifiable*: The users should be able to use their OS knowledge to guess the abstractions and the functions they perform. For this, the subsystems of the OS should be directly identifiable in the wrappers. This makes learning and understanding easier. To achieve this goal, we have created abstractions such as `MemoryManager`, `ProcessManager` etc., which directly correspond to memory management and process subsystems of Linux kernel. It is sufficient for the users to understand these wrappers to perform most of the tasks with the subsystems.
2. *Facilitate housekeeping within the wrappers*: Often, for an OS, a function should be called only after performing certain operations such as locking. Keeping track of locking and unlocking makes the programmer's task difficult. Therefore, it would be helpful if most of the

housekeeping is done within the wrappers so that the programmer can concentrate on the primary task.

3. *Reusability*: Reusability is an important characteristic of OO design. This requires the wrappers to be simple, so that the concerns are separated well. Our abstractions are made highly cohesive by localizing the interactions. Hence, defining new wrappers that can extend existing ones becomes easier.
4. *Performance overhead of wrappers should be less*: Performance is one of the most important qualities required by the applications such as kernel modules. When wrappers are used in the development of kernel modules, they introduce an extra level of indirection, i.e. an extra function call to do the same function, which creates a performance overhead. The wrapper system will not be useful if the performance overhead due to indirection is high.

## 4.2. Design of wrapper-based kernel

Linux is a large monolithic kernel, where in different subsystems such as memory management, process scheduling, inter process communication and device drivers run in a single executable file. The kernel provides system call interface to the applications for requesting various services of the kernel.

Linux provides support for loading and linking modules to the kernel at runtime. This loadable module support removes the need for recompiling the entire kernel for supporting additional modules. Device drivers and file systems are examples of loadable modules.

Although the kernel is monolithic, conceptually the kernel can be viewed as two parts, the core kernel, containing different subsystems such as memory management, process scheduling, inter-process communication, interrupt handling, etc., and the other part the loadable modules, such as device drivers.

The Linux kernel handles the device drivers in the following way. In Linux, the devices are considered as files. Hence, when a user application wishes to read or write to a device, it issues a read/write system call on the corresponding device file. The virtual file system (VFS) routes the system call to an appropriate device function. The VFS provides callbacks for operations of device drivers. The device drivers define the callbacks and register them with the VFS. As every device is a kind of file, the interface of the callbacks for device drivers is provided by the VFS. For example, the *file\_operations* structure that provides the interface for character drivers is defined in *fs.h*, which belongs to VFS. Hence, in the Linux kernel, the core kernel communicates with device drivers through the VFS, whereas the device drivers communicate directly with the core kernel.

We provide wrappers for the core kernel of Linux. The design of our wrapper-based kernel is shown in Figure 4.

In the wrapper-based kernel, it is not possible to register the C++ callbacks with C kernel. Hence, in our wrapper library, we define the C interface for callbacks of the kernel, which in turn calls the polymorphic driver methods of wrappers and registers it with the VFS. Hence, when the kernel wishes to communicate with a driver, it makes a call to the VFS, which in turn forwards it to wrappers. The wrapper does the processing and forwards it to appropriate driver. Hence, in the wrapper-based kernel, the communication from kernel to device drivers is through the VFS and wrappers.

However, when the driver wishes to communicate with the core kernel, it makes calls to appropriate wrapper methods, which in turn delegates the request to the underlying Linux kernel. Hence,

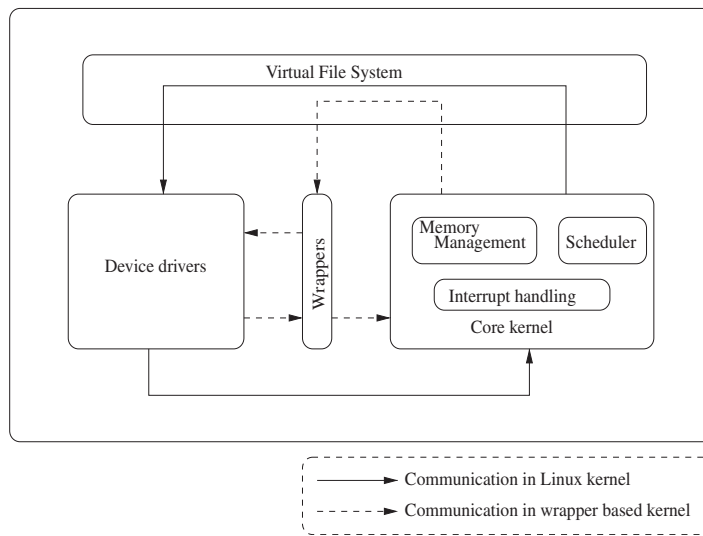


Figure 4. Design of wrapper-based kernel.

the communication of drivers to the core kernel is through wrappers. This feature can further be exploited for storing the state information, which can be used for shielding the device driver failure from kernel.

The design of wrappers to core kernel is based on the Facade pattern as discussed in [12,13]. The intent of the pattern is to ‘provide a uniform interface to a set of interfaces in the system. Facade defines a higher level interface which makes the subsystems easier to use’ [12]. In simple words, it provides a higher level interface by abstracting the details that hamper the understanding of the interface. There is no difference to the client in terms of functionality, but it only enhances the ease of use.

The kernel is abstracted into the following three major subsystems—memory, process and interrupt. These three subsystems are necessary for the device driver programmers. Some of the important wrappers and the corresponding services are listed in Table I.

Memory management subsystem deals with the dynamic memory allocation for kernel usage. KernelMemoryAllocator wraps the functions that use `kmem_cache` data structure. This abstraction can be used to allocate memory for kernel data structures. The functionality provided by this abstraction is analogous to `malloc`. MemoryManager wrapper provides the functionalities required to manage the memory of a process. NonContiguousMemory and Pager wrappers provide functionalities for dealing with virtual memory.

The Scheduler and Waitqueue wrappers provide the functionality for managing a process. The Waitqueue abstraction is useful for character drivers.

Timer and Interrupt management subsystems handle the timing measurement and interrupt-handling tasks, respectively. InterruptHandler wrapper facilitates working with IRQ data structures, which is an important wrapper for all the device drivers. It contains many useful functions such as `request_irq`, `save_irq`, `can_request_irq`, `probe_irq_on`, `probe_irq_off`, `enable_irq` and `disable_irq`. The irq handlers are defined as strategies (class Handler) whose instance is present in InterruptHandler

Table I. Wrappers for Linux kernel.

S.no.	Subsystem	Wrapper	Wrapper functions
1	Memory	KernelMemoryAllocator	kmem_cache_create, kmem_cache_reap, kmem_cache_shrink, kmem_cache_free
2	Memory	MemoryManager	mm_alloc, exit_mmap
3	Memory	VirtualMemoryManager	find_vma, vma_merge
4	Memory	NonContiguousMemory	vmalloc, vfree
5	Memory	Pager	alloc_pages, get_free_pages, free_pages, free_pages_bulk
8	Process	Waitqueue	init_waitqueue_head, init_waitqueue_entry, add_wait_queue, remove_wait_queue
9	Process	Scheduler	schedule_timeout_uninterruptible, schedule_timeout_interruptible
10	Interrupt	InterruptHandler	enable_irq, free_irq, disable_irq, local_irq_save, local_irq_restore, local_irq_enable, request_irq
11	Interrupt	Tasklet	tasklet_init, tasklet_disable_nosync, tasklet_disable, tasklet_enable
12	Timer	TimerManager	do_gettimeofday, update_times

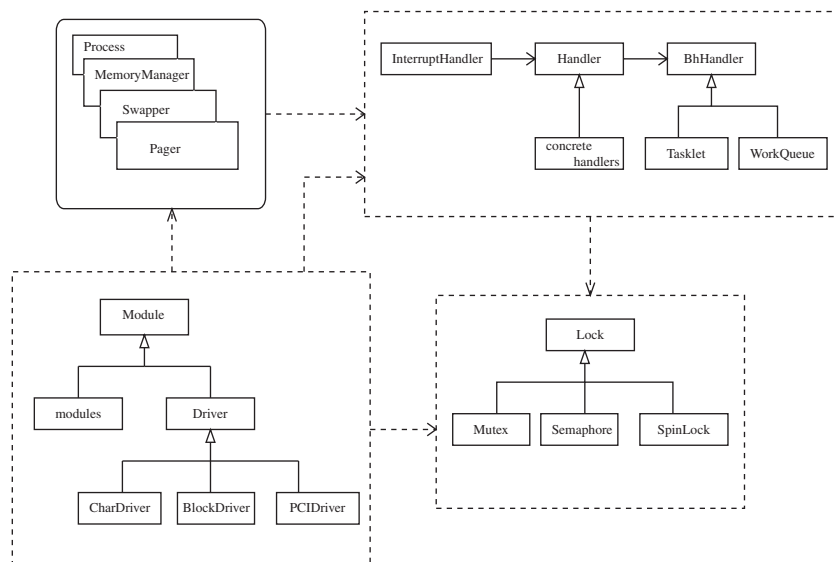


Figure 5. Class diagram of wrappers.

class. A static handler function is statically registered in the wrapper framework. This handler function internally calls the polymorphic handler function of `InterruptHandler`, which can be defined by the user.

A simple class diagram of wrappers is shown in Figure 5.



## 5. IMPLEMENTATION

In this section we discuss the implementation of C++ wrappers for the Linux kernel.

A patch that gives the runtime support of C++ in Linux kernel is provided in [14]. The patch provides all the basic C++ features, such as virtual functions, exceptions and inheritance, in the Linux kernel with a minimal overhead. It reduces the developer's unnecessary effort in masking the variables in the Linux kernel that have names of C++ keywords such as `new`, `delete`, etc. If the user wishes to use the kernel application programming interface (API) in C++ modules, the C header files should be enclosed between `begin_include.h` and `end_include.h`, which are provided by the patch. However, it is not a good practice to mix C and C++ codes. Hence, instead of using the kernel API directly, the C++ wrappers can be used.

The wrappers have been provided as include files in `'include/oowrap'`. The wrappers are organized in a similar way to that of Linux source code organization. For example, the memory-related wrappers are placed in `'oowrap/mm'`, and process-related wrappers are placed in `'oowrap/process'`. The process-related functions are not placed separately in Linux kernel directory tree. The reason behind allocating a separate directory for process-related wrappers is for the clear realization of the subsystems.

A simple wrapper `'InterruptHandler'` for interrupt handling is shown in Figure 6. This is a wrapper around interrupt-handling functions `enable_irq`, `disable_irq`, `probe_irq_on` and `request_irq`.

If a developer wishes to allocate memory for his/her kernel data structures, he/she creates an object of `KernelMemoryAllocator` and calls the `alloc` function in it, which forwards the call to the underlying kernel's `kmalloc` function.

The source code of Linux kernel along with wrappers is available for download in [15].

```
#include <begin_include.h>
#include <linux/slab.h>
#include <end_include.h>

class InterruptHandler
{
public:
    void m_enableIrq(unsigned int irq)
    {
        enable_irq(irq);
    }
    void m_disableIrq(unsigned int irq)
    {
        disable_irq(irq);
    }

    void m_freeIrq(unsigned int irq, void* dev_id)
    {
        free_irq(irq, dev_id);
    }
};
```

Figure 6. Implementation of wrappers.

## 6. HOW TO USE THE WRAPPERS

The wrappers can be used for programming modules in an OO fashion. It provides the programmer with the flexibility of programming completely in OO without mixing the procedural and OO paradigms. A simple example in Figure 7 shows the flexibility of using C++ in kernel modules. This module does not give the complete code. It just shows how modules such as device drivers can benefit by using C++ wrappers.

This example shows a simple driver program. Many vendors produce different drivers for a device. The applications may use any of these drivers. The driver class defines the interface that different drivers should confirm to. ConcreteDriver1 and ConcreteDriver2 implement this interface. Applications can use any of these drivers by setting the reference of driver to appropriate ConcreteDriver. *Terminal* is a wrapper that provides functions for printing onto the terminal.

The example shown is a polymorphic driver that can be switched at runtime based on the requirements of the user. The user has to write a driver that extends the *Driver* interface and define the abstract functions of *Driver*. Based on the concrete object the user assigns to the abstract *Driver* reference, the corresponding driver will be called.

To develop the same polymorphic driver in C, it needs considerably more development and testing effort than in C++. This is because the developer should first implement a virtual table to support driver switching and handle operations on this virtual table carefully while writing driver programs. This includes a lot of development and maintenance work. Also it is more prone to errors than the

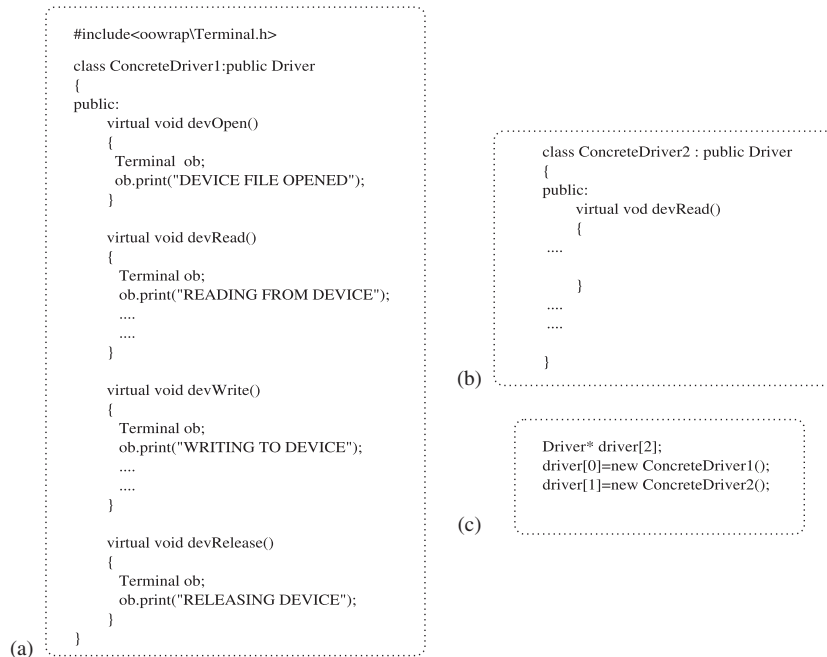


Figure 7. Usage of wrappers.

C++ version as it should be done manually, whereas it is handled by the compiler if developed in C++.

## 7. CASE STUDY: NETWORK DEVICE DRIVER

Device drivers are one of the main components that interact with the core kernel. In Linux, device drivers can be inserted at runtime as modules or can be compiled statically into the kernel.

We have re-engineered two network device drivers called `ne2k_pci.c` and `8210.c` from C to C++. The `ne2k_pci` is a device driver for PCI NE2000 clones, and `8210.c` is the driver for 8210 series Ethernet cards. These drivers are used frequently when the network connection of the system is used, enabling us to estimate performance degradation caused due to wrappers.

The functions of the driver are grouped into a class structure, and wherever the kernel interface is used, we replaced them with the instances of the appropriate C++ wrappers. A small part of the network driver code is shown in Figure 8.

The re-engineered driver uses two wrappers, the `InterruptHandler` and `Terminal`, which are related to interrupt handling and printing-related functions. The driver was completely written in C++, without using any of the kernel interfaces directly.

### 7.1. Evaluation of case study

Apart from being an OO interface to the Linux kernel, we identified certain qualities (in Section 4.1) in the domain of OSs, which improve the usability of the wrappers. Among these qualities, structure identifiability, housekeeping within wrappers and reusability affect the development and maintainability effort of kernel modules. In this section, we explain how these qualities are achieved using wrapper-based design. As performance is considered the most important quality factor for kernel modules apart from the other three, we provided detailed analysis of it in Section 7.2.

1. *Structure identifiability*: During the development of kernel modules, it becomes easy for the developer to use abstractions that represent the concepts of the domain (OS). Hence, the kernel is

```
class ne2k_pci_OO
{
private:
    Terminal obj;
    InterruptHandler interrupt_handler_obj;
public:

    int ne2k_pci_open();
    int ne2k_pci_close();
    void ne2k_pci_block_input();
    void ne2k_pci_block_output();
    .....
    .....
    void ne2k_pci_get_drvinfo();
    int __devinit ne2k_pci_init_one ();

}ne2k_pci_obj;
```

Figure 8. A sample network driver reengineered to C++.

abstracted into a set of high-level wrappers, which directly correspond to different subsystems of the kernel. The developer just needs to learn a set of wrappers of the kernel to develop a module. This decreases the learning time, which in turn reduces the development time. For example, to develop the network driver shown in Figure 8, all that the developer needs to know about the kernel is the InterruptHandler (for handling interrupts) and Terminal (for printing driver debug information) wrappers.

2. *Housekeeping within wrappers*: During kernel module development, care should be taken to ensure that certain aspects such as right parameters are being passed to functions; the locking- and unlocking-related functions are called correctly before and after a specific function call and a lock has been achieved before using a data item. These aspects are repetitive and failing to do so may sometimes cause a severe corruption to kernel data structures and may even lead to a kernel crash. Ensuring these aspects during driver development complicates the developers' work. Hence, encapsulating them in the wrappers reduces considerable development effort.

For example, in the network device driver, the parameter 'hdr' is identified as a critical parameter and its incorrect usage might lead to a kernel crash [16]. Care should be taken to pass correct value to this parameter. A housekeeping wrapper, called `ne2kpci.filter`, was associated with the driver, which automatically checks the usage of this parameter and restricts improper usage. Similarly, locks can be obtained by the use of scoped locking pattern [17] in wrappers.

3. *Reusability*: The device driver hierarchy provides blackbox reuse, where in order to develop a new driver, the user just needs to extend it from an appropriate class of driver and define the specific parts of that driver. The reuse can also be achieved by abstracting the architecture-specific parts of code and composing the driver with specific architecture object. In either case, the reuse achieved is considered better than the reuse achieved in the procedural paradigm.

## 7.2. Performance evaluation

We evaluated the performance of C driver in the original Linux kernel and C++ enabled Linux kernel. We also measured the performance of C++ driver in C++ enabled Linux kernel.

The performance of the C and C++ drivers is observed under both idle and load conditions. The performance is measured using the *netperf* tool [18], which is a benchmarking tool for network applications. The tool performs various networking tests by changing the input data size and provides results in terms of transaction rate (a transaction being one send and receive).

As expected, there is a degradation in the performance in the C++ driver. The degradation is as little as 2%. Given the advantages of writing the device driver in C++, the overhead of 2% in terms of performance is minimal.

The graphs corresponding to TCP and UDP tests performed by *netperf* tool are shown in Figures 9 and 10.

Further experimentation was carried out using *Hbench* microbenchmarking [19] tool to compare the performance of procedural and OO network driver on the modified Linux kernel. *Hbench* is a collection of benchmarking tools, which measure various aspects of OS such as memory read, write, network bandwidths and latencies. The readings are taken over 10 iterations and the median of the values is presented. We have limited our experimentation to the tests corresponding to the




Figure 9. Driver performance for TCP.

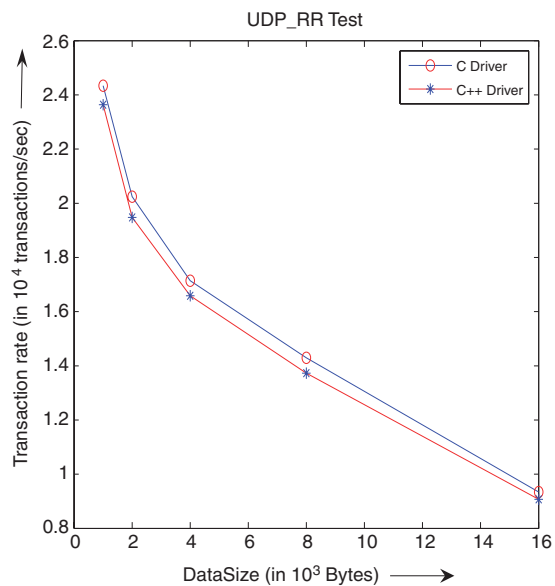


Figure 10. Driver performance for UDP.

network, as the main objective was to check the performance of network driver. The tests performed using *Hbench* tool are:

1. TCP bandwidth test (bw\_tcp): This test measures the bandwidth attainable when transferring data through a TCP connection between two processes. This test is performed using both C and C++ drivers, varying the buffer size being transmitted between the processes. The test outputs the attainable bandwidth in mbps; the greater the value, the better the performance. A slight degradation in the performance is observed with C++ driver as shown in Figure 11.
2. TCP connection latency test (lat\_connect): This test measures the latency in establishing a TCP connection between the two hosts. It outputs the latency in microseconds. The lesser the value, the better the performance.
3. TCP transaction latency test (lat\_tcp): This test measures the latency of a 1-byte ping-pong between two processes connected via a TCP connection.
4. UDP latency test (lat\_udp): This test measures the latency of a 1-byte ping-pong between two processes using UDP datagrams to transfer the data.
5. RPC latency (lat\_rpc): This test measures the latency of RPC calls via UDP and TCP transports.

The performance results collected using *Hbench* tool are shown in Figures 11 and 12. In addition to a slight degradation in the performance, we have observed a slight increase in the size of the wrapper-based C++ kernel when compared to the original kernel. An increase of 118 kB is observed in the modified kernel image size. This change is attributed to the addition of C++ libraries to the kernel.




Figure 11. TCP bandwidth tests of Hbench tool.

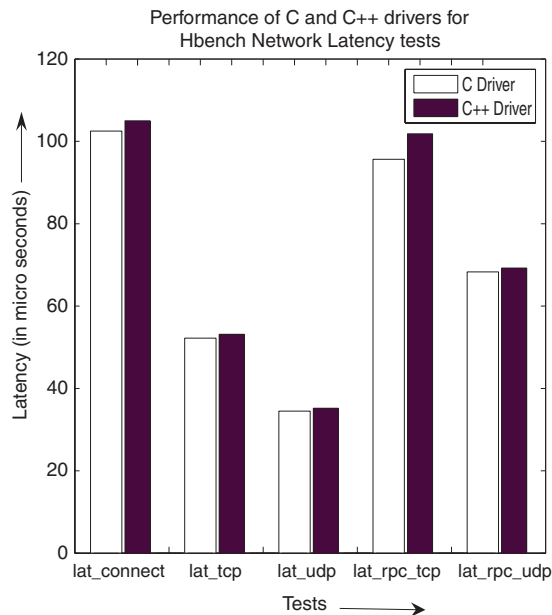


Figure 12. Network latency tests of Hbench tool.

## 8. RELATED WORK

### 8.1. OO operating systems

Choices is one of the most popular OO OS developed in C++ [6]. Choices encourages reuse by including mechanisms for inheritance and polymorphism. It provides a framework for reusing the design within the OS.

Spring is another OO OS, which uses interface definition language for specification of system interfaces [20]. It separates the interface from implementation so that both can vary independently.

### 8.2. OO wrappers for OSs

Wrappers have been provided for Mach OS [21]. Mach is a microkernel-based OS in which only minimal components are kept in the kernel space. All the remaining subsystems run in the user space. Hence, this includes a lot of communication between the subsystems that run in the user space and the kernel. The main objective of providing OO wrappers for Mach kernel is to hide state management issues involved in communicating with microkernel and to make the communication with the kernel easier [21].

Our work is different from that of Mach work in the sense that we have provided wrappers in the kernel space for the in-kernel API for the benefit of kernel modules such as device drivers. Kernel modules can access the kernel directly either through kernel API or through wrapper interface. In this way, the existing modules need not be changed and the new modules can be developed using

the wrapper interface using C++. The wrapper functions have been made inline wherever possible, so as to reduce the performance overhead.

### 8.3. I/O kit

I/O kit is a device driver model specially built to support the MACH kernel of Mac OS X. It is an OO framework built in C++ with a layered architecture. The main architectural features of I/O kit include hardware modeling and layering of driver objects, driver matching, I/O kit registry and I/O catalog, and I/O class hierarchy [22].

Although I/O kit is designed specifically for device drivers, the wrappers we provide can be used to build any of the external kernel modules for Linux. I/O kit handles the requests to a driver by processing it at different layers (which is developed in C++), whereas in our wrappers we delegate it to the Linux kernel (which is developed in C). Hence, we expect the performance of device drivers to be better in our case.

## 9. CONCLUSION AND FUTURE WORK

C++ wrappers have been built for the core kernel of Linux OS and it can be used by kernel module programmers for developing their programs in C++, which enhances maintainability and reuse for their programs. A number of device drivers have been re-engineered to C++ and a detailed performance study has been presented.

The wrapper support for Linux kernel opens up new options of intercepting the messages that go to the kernel. We are exploiting this option to provide a new security mechanism based on biometric devices. We are also studying the usability of filter objects [23] in Linux kernel to enhance the security of the kernel. Using wrappers and filter object model, it is possible to build interception of messages to kernel. This approach can provide device driver isolation as provided by NOOKS [24]. The wrappers for Linux kernel can be downloaded at <http://dos.iitm.ac.in/MOOL/>.

## ACKNOWLEDGEMENTS

The authors acknowledge the Department of Information Technology, Government of India for providing financial support for the Linux redesign project. The project review team has provided critical reviews to the project.

## REFERENCES

1. Yu L, Schach SR, Chen K, Offutt J. Categorization of common coupling and its application to the maintainability of the Linux kernel. *IEEE Transactions on Software Engineering* 2004; **30**(10):694–706.
2. Yu L, Schach SR, Chen K, Heller GZ, Offutt J. Maintainability of the kernels of open-source operating systems: A comparison of Linux with FreeBSD, NetBSD, and OpenBSD. *The Journal of Systems & Software* 2006; **79**(6):807–815.
3. Reddy VK, Janakiram D. Cohesion analysis in Linux kernel. *Proceedings of the XIII Asia Pacific Software Engineering Conference*. IEEE Computer Society: Washington, DC, 2006; 461–466.
4. The Linux-kernel Mailing List FAQ. <http://www.kernel.org/pub/linux/docs/lkml/> [3 August 2007].
5. Wheeler DA. More than a Gigabuck: Estimating GNU/Linux size. *Sixth International Workshop on Program Comprehension*, 1998.
6. Campbell RH, Islam N, Raila D, Madany P. Designing and implementing choices: An object-oriented system in C++. *Communications of the ACM* 1993; **36**(9):117–126.



7. Yokote Y. The Apertos reflective operating system: The concept and its implementation. *ACM SIGPLAN Notices* 1992; **27**(10):414–434.
8. Stroustrup B. *The Design and Evolution of C++*. ACM Press/Addison-Wesley Publishing Co.: New York/Reading, MA, 1995.
9. Liu SS, Wilde N. Identifying objects in a conventional procedural language: An example of data design recovery. *Conference on Software Maintenance*, 1990; 266–271.
10. Canfora G, Cimitile A, Munro M. An improved algorithm for identifying objects in code. *Software—Practice and Experience* 1996; **26**(1):25–48.
11. Gall H, Klösch R. Finding objects in procedural programs: An alternative approach. *Proceedings of the Second Working Conference on Reverse Engineering*. IEEE Computer Society: Los Alamitos, CA, 1995; 208–216.
12. Gamma E, Helm R, Johnson R, Vlissides J. *Design Patterns: Elements of Reusable Object-oriented Software*. Addison-Wesley Longman Publishing Co.: Reading, MA, 1995.
13. Schmidt D. Wrapper facade: A structural pattern for encapsulating functions within classes. *C++ Report*, vol. 11(2), 1999; 40–50.
14. C++ in Linux Kernel. <http://netlab.ru.is/exception/linuxcxx.shtml> [3 August 2007].
15. Minimalistic Object Oriented Linux Homepage. <http://dos.iitm.ac.in/MOOL> [3 August 2007].
16. Pendyala N. Implementation of message filters in Minimalistic Object Oriented Linux. *MTech Thesis*, <http://dos.iitm.ac.in/MOOL>.
17. Schmidt DC, Rohnert H, Stal M, Schultz D. *Pattern-oriented Software Architecture: Patterns for Concurrent and Networked Objects*. Wiley: New York, U.S.A., 2000.
18. Netperf Home Page. <http://www.netperf.org> [3 August 2007].
19. Brown A, Seltzer M. Operating system benchmarking in the wake of Imbench: A case study of the performance of NetBSD on the Intel x86 architecture. *Proceedings of the 1997 ACM SIGMETRICS International Conference on Measurement and Modeling of Computer Systems*, 1997; 214–224.
20. Islam N, Campbell RH. Latest developments in operating systems. *Communications of the ACM* 1996; **39**(9):38–40.
21. Kurtzman S, Dattatri K. Design goals of object-oriented wrappers for the Mach micro kernel. *Compcon'95*. IEEE Computer Society: Washington, DC, 1995; 367–372.
22. I/O Kit Fundamentals. <http://developer.apple.com/DOCUMENTATION/DeviceDrivers/Conceptual/IOKitFundamentals/> [3 August 2007].
23. Joshi RK, Vivekananda N, Ram DJ. Message filters for object-oriented systems. *Software—Practice and Experience* 1997; **27**(6):677–699.
24. Swift MM, Annamalai M, Bershad BN, Levy HM. Recovering device drivers. *Proceedings of the 6th USENIX Symposium on Operating Systems Design and Implementation*, 2004.